

I8. DATA LAKE VS. DATA WAREHOUSE



ADS ISE II

-PARSHWA HERWADE (22510064)

Data Lake vs. data warehouse

- Introduction & Context
- Definition and Purpose
- Architectural Overview
- Key Differences & Decision Factors
- Pros, Cons, and Use Cases
- Emerging Trends & Hybrid Approaches
- Summary & Conclusion
- References & Further Reading

Literature Survey

- Data Quality & Structure:

Warehouses: Enforce clean, structured data via rigorous ETL.

Lakes: Ingest raw data with flexible schemas, requiring robust metadata management.

- Query Performance:

Warehouses: Optimized with pre-aggregated data and indexing for rapid, predictable queries.

Lakes: Rely on distributed processing (e.g., Apache Spark) to handle complex queries over diverse data.

- Scalability & Flexibility:

Warehouses: Effective in regulated environments but can be resource-intensive to scale.

Lakes: Offer scalable, cost-effective storage for large, varied data sets.

- Industry Adoption:

Warehouses are well-established in traditional BI environments.

Lakes and emerging hybrid models (Lakehouses) are validated by recent research in big data analytics.



Introduction & Context

- Background:


1. Rapid data growth from IoT, social media, and enterprise applications.
2. Increasing need to manage structured and unstructured data.

- Why Compare?

3. Organizations must choose the right data storage/processing architecture based on business needs.
4. This comparison helps decision-makers balance cost, scalability, and performance.

- Visual/Statistic:

5. Consider an infographic showing data growth trends globally.



What is a Data Warehouse?

- Definition:

1. A centralized repository designed to store structured data for reporting and analysis.

- Purpose and Usage:

2. Facilitates business intelligence, historical analysis, and decision-support.

- Core Features:

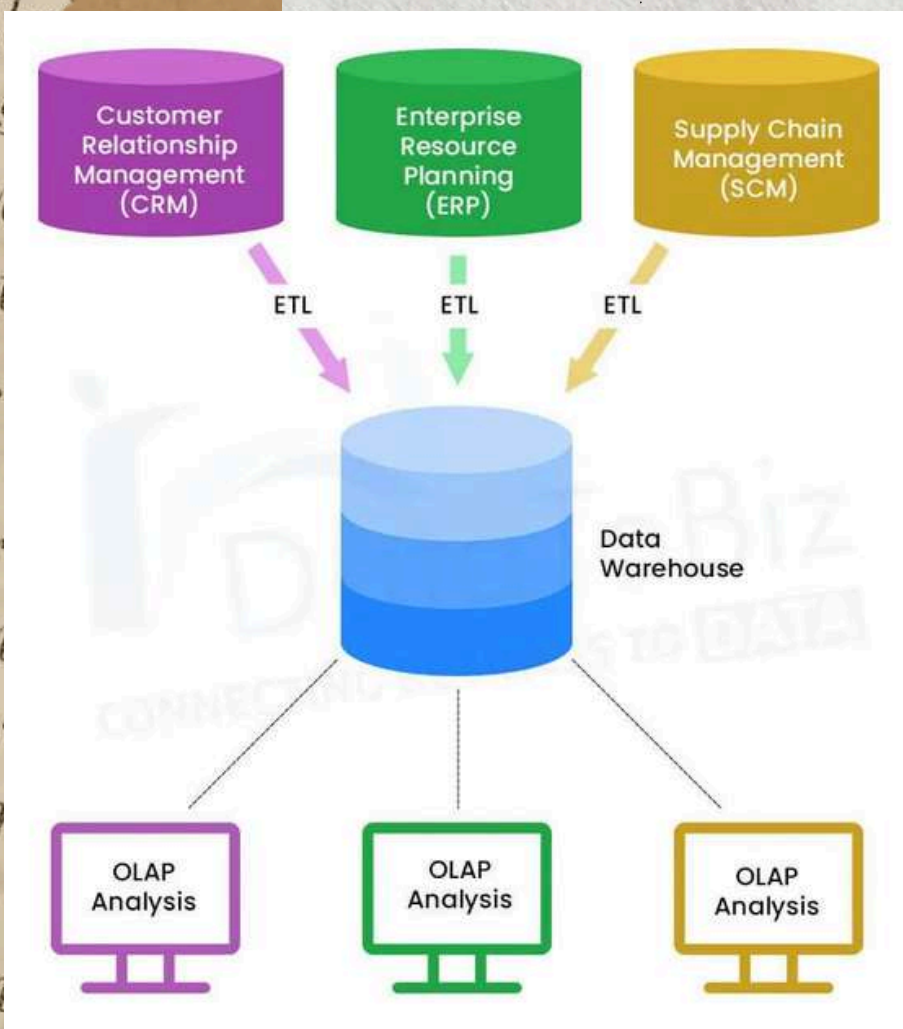
3. Schema-on-Write: Data is cleansed, transformed, and structured before storage.

4. Optimized for OLAP: Supports fast querying and reporting.

5. Data Integration: Combines data from disparate sources through ETL/ELT processes.

- Visual Aid:

6. Diagram of ETL flow from source systems to Data Warehouse beside.



What is a Data Lake?

1. Definition:

- a. A storage repository that holds raw, unstructured, semi-structured, and structured data in its native format.

2. Purpose and Usage:

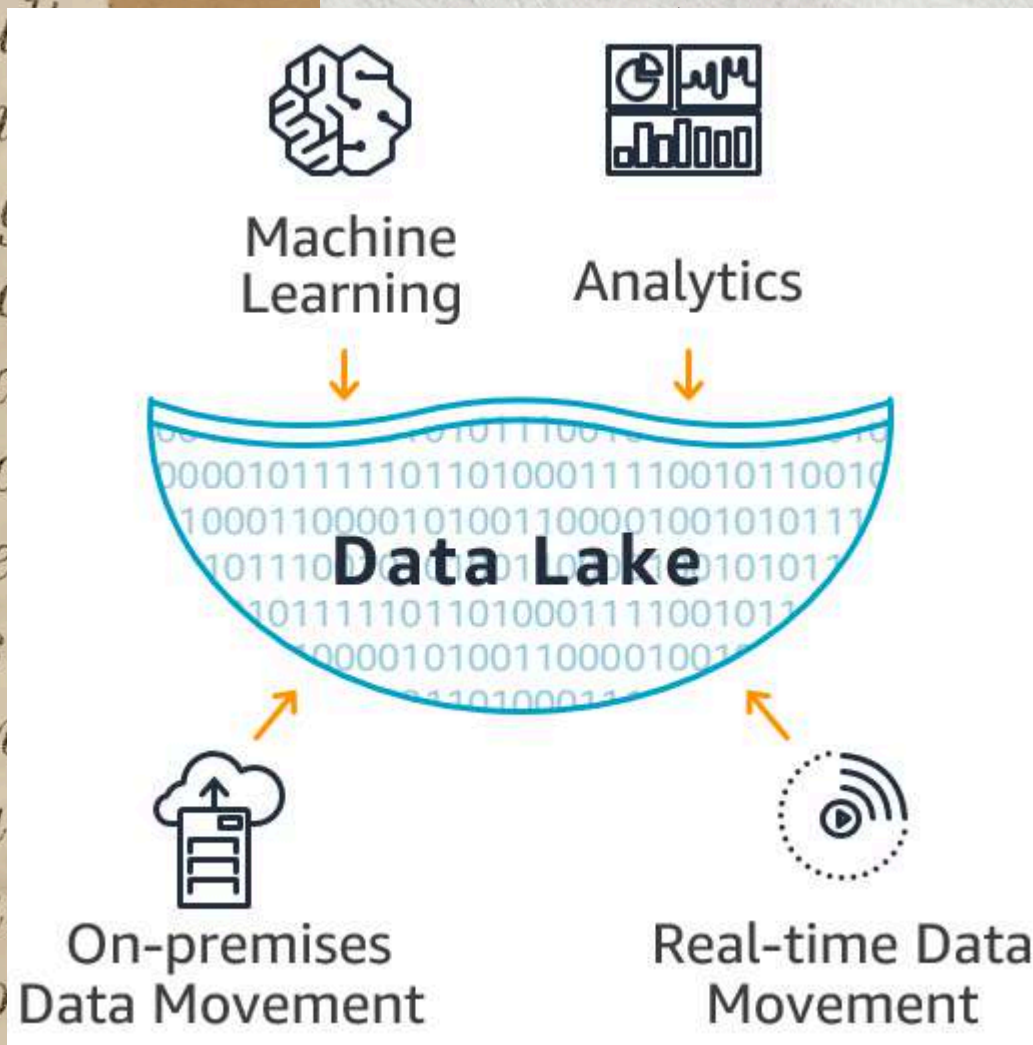
- a. Enables flexible data storage; supports data exploration, analytics, and machine learning.

3. Core Features:

- a. Schema-on-Read: Structure is applied when data is accessed.
- b. High Scalability: Uses distributed storage systems (e.g., Hadoop, cloud object storage).
- c. Data Diversity: Supports video, images, logs, and text as well as traditional structured data.

4. Visual Aid:

- a. Diagram with a layered approach: ingestion → raw data storage → processing framework (e.g., Apache Spark).



Architectural Comparison & Data Flow

1. Data Warehouse Architecture:

- ETL Process: Extract, Transform, Load data with a focus on data quality and consistency.
- Storage: Relational databases on dedicated hardware.
- Query Optimization: Uses indexing and aggregations to support complex queries.

2. Data Lake Architecture:

- Data Ingestion: Stores data in raw form with minimal upfront processing.
- Storage: Utilizes distributed file systems or cloud object storage.
- Processing Frameworks: Analysis through scalable systems (Apache Spark, Presto, etc.).
- Governance Concerns: Emphasis on metadata management to avoid “data swamp” issues.

3. Visual Comparison:

- ETL vs. raw data ingestion, schema-on-write vs. schema-on-read, cost models, and performance aspects.



Key Differences & Decision Factors

1. Organization & Schema Approach:

- Data Warehouse: Predefined, rigid schema for consistent data.
- Data Lake: Flexible, varied schema applied at read time.

2. Performance and Efficiency:

- Warehouse: Faster queries on structured data, reliable performance.
- Lake: Potential processing delays; requires robust indexing/metadata for efficient querying.

3. Cost & Scalability:

- Warehouse: Higher cost and maintenance overhead due to specialized hardware and ETL.
- Lake: Lower storage cost with scalable, commodity hardware; hidden costs in processing.

4. Security & Governance:

- Warehouse: Mature data governance and security frameworks.
- Lake: Demands additional data cataloging and governance tools.

5. Decision Factors:

- Evaluate your use case: reporting vs. exploratory analytics, structured vs. diverse data sources.

Pros and Cons

PROS

CONS

• I. Data Warehouse Advantages:

- a. High data quality and consistency, optimized performance for complex queries.
- b. Strong data governance, which is ideal for regulatory compliance.

2. Data Warehouse Limitations:

- a. Inflexible structure that may struggle with rapidly changing data types.
- b. Higher cost and slower integration of new/unstructured data sources.

3. Data Lake Advantages:

- a. Flexibility to store all types of data with lower upfront processing costs.
- b. Scalability ideal for big data analytics and machine learning applications.

4. Data Lake Limitations:

- a. Can degrade into a “data swamp” without proper metadata management.
- b. Requires more complex processing and security measures for efficient querying.

5. Visual Comparison:

- a. A bullet list or two-column table to directly compare benefits and drawbacks.

Use Cases & Industry Examples

I. Data Warehouse Use Cases:

- a. Enterprise reporting, sales analysis, financial forecasting, and historical trend analysis.
- b. Examples: Financial institutions, retail chains using centralized data for decision support.

2. Data Lake Use Cases:


- a. Data science projects, real-time analytics, storing sensor, social media, and IoT data.
- b. Examples: E-commerce companies analyzing user behavior or healthcare institutions storing medical imaging alongside clinical data.

3. Visual Examples:

- a. Include icons/logos of industries, or sample case study snapshots.

4. Real-World Insights:

- a. Mention how hybrid approaches have been successfully used in industries like finance and healthcare.



Emerging Trends & Hybrid Approaches

I. Trends in Data Management:

- Hybrid Solutions: Adoption of Data Lakehouse architectures that bring together low-cost storage and robust query performance.
- Cloud Integration: Increasing move towards integrated cloud platforms (AWS, Azure, Google Cloud) that support both data lakes and warehouses.
- Advancements in Tools: Introduction of technologies like Delta Lake and Apache Iceberg to improve metadata management and query performance on data lakes.

2. Future Outlook:

- Evolution of real-time analytics with machine learning integration.
- Increasing emphasis on building flexible yet governed data environments.



References

1. Books & Publications:

- “The Data Warehouse Toolkit” by Ralph Kimball.
- “Data Lake Architecture: Designing the Data Lake and Avoiding the Garbage Dump” (various authors).


2. Research Papers:

- Papers on data lake governance and hybrid data architectures from IEEE, ACM, or similar reputable sources.

3. Online Resources:

- Articles from cloud providers (AWS, Azure, Google Cloud) comparing data lakes and data warehouses.
- Technical blogs and white papers such as “Understanding the Data Lakehouse”.





Thank you

The background features a collage of torn, aged paper. The paper has a yellowish-brown hue and irregular, deckled edges. In the top left corner, there is a black line drawing of a large, stylized letter 'Q'. The paper is layered, with some pieces showing Latin text in a serif font. The text is partially obscured by the torn edges and the central 'Thank you' text. In the bottom right corner, the spiral binding of a notebook is visible, with the metal or plastic rings holding the pages together. The overall composition suggests a vintage or handmade aesthetic.